

# Section overview

## Lets go Markov

### Policy

- 1 Motivation
- 2 Reinforcement Learning 101
- 3 Lets go Markov
  - Markov Process
  - Markov Reward Process
  - From state to action:
- 4 Markov Decision Process
- 5 Dynamic Programming
- 6 Model-Free Learning
- 7 Model-Free Control

# Markov Process

## Definition

A Markov Process is a tuple  $(\mathcal{S}, \mathcal{P})$ ,

$\mathcal{S}$  is a set of states

$\mathcal{P}_{ss'}$  is a state transition probability matrix,

$$\mathcal{P}_{ss'} = P[S_{t+1} = s' | S_t = s] \quad (2)$$

A **Markov process** generates a chain of Markov states governed by probabilistic transitions.

# Markov Property


## Definition

A state  $s_t$  is **Markov** if and only if  $P[s_{t+1}|s_t] = P[s_{t+1}|s_1, \dots, s_t]$

The future is independent of the past given the present.

- The present state  $s_t$  captures all information in the history of the agent's events.
- Once the state is known, then any data of the history is no longer needed.

# Example: University life Markov Process



Round states (transient states), Square states (**terminal states**).

2

# State Transition Probabilities

For a Markov state  $s$  and successor state  $s'$ , the state transition probability is defined by  $\mathcal{P}_{ss'} = P[s_{t+1} = s' | s_t = s]$ . The state transition matrix defines transition probabilities from all states  $s$  to all successor states  $s'$ .

Because all transition probabilities have to be accounted for to give a total probability of 1, we have  $\sum_{s'} \mathcal{P}_{ss'} = 1$  (we need to end up somewhere after leaving  $s$ , including returning to  $s$ ). We choose the matrix row notation so that the rows have to sum to one.

## Definition (Stationarity)

If the  $P[s_{t+1} | s_t]$  do not depend on  $t$ , but only on the origin and destination states, we say the Markov chain is **stationary** or **homogenous**.

# Markov Reward Process

A **Markov Reward Process** (MRP) is a Markov chain which emits rewards.

## Definition (Markov Reward Process)

A Markov Reward Process is a tuple  $(\mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma)$


$\mathcal{S}$  is a set of states

$\mathcal{P}_{ss'}$  is a state transition probability matrix

$\mathcal{R}_s = \mathbb{E}[r_{t+1} | S_t = s]$  is an expected immediate reward that we collect upon departing state  $s$ , this reward collection occurs at time step  $t + 1$

$\gamma \in [0, 1]$  is a discount factor.

# Example: University life Markov Reward Process



What do samples from this process look like?

# Return

## Definition (Return)

The **return**  $R_t$  is the total discounted reward from time-step  $t$ .

$$R_t = r_{t+1} + \gamma r_{t+2} + \cdots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \quad (3)$$

The factor  $\gamma \in [0, 1]$  is how we **discount** the present value of future rewards.

The value of receiving reward  $r$  after  $k + 1$  time-steps is  $\gamma^k r$ .

The discount values immediate reward higher than delayed reward:

- $\gamma$  close to 0 leads to "**myopic**" (short-sighted) evaluation.
- $\gamma$  close to 1 leads to "far-sighted" evaluation.



# University life MRP returns

Sample returns for Student MRP:

Starting from  $S_1 = C1$  with  $\gamma = \frac{1}{2}$

$$R_1 = r_2 + \gamma r_3 + \dots + \gamma^{T-2} r_T$$

$T$  is for the time it takes to reach the terminal state.

## Example (Sample Runs)

C1 C2 C3 Pass Sleep	$R_1 = -2 + \frac{1}{2} \times -2 + \frac{1}{2}^2 \times -2 + \frac{1}{2}^3 \times 10$
C1 FB FB C1 C2 Sleep	$R_1 =$
...	$-2 + \frac{1}{2} \times -1 + \frac{1}{2}^2 \times -1 + \frac{1}{2}^3 \times -2 + \frac{1}{2}^4 \times -2$
C1 FB FB FB ...	$R_1 = -2 + \frac{1}{2} \times -1 + \frac{1}{2}^2 \times -1 + \frac{1}{2}^3 \times -1 + \dots$

What is the value of being in C1, C2, C3?

## Why discounting is a good idea?

Most Markov reward processes are discounted with a  $\gamma < 1$ . Why?

- Mathematically convenient to discount rewards
- Avoids infinite returns in cyclic or infinite processes
- Uncertainty about the future may not be fully represented
- If the reward is financial, immediate rewards may earn more interest than delayed rewards
- Human and animal decision making shows preference for immediate reward
- It is sometimes useful to adopt undiscounted processes (i.e.  $\gamma = 1$ ), e.g. if all sequences terminate and also when sequences are equally long (why?).


# State Value Function

## Definition (State value function)

The state value function  $v(s)$  of an MRP is the **expected return**  $R$  starting from state  $s$  at time  $t$ .

$$v(s) = \mathbb{E}[R_t | S_t = s] \quad (4)$$

## Example (Golf)



# Bellman Equation for MRPs

$$v(s) = \mathbb{E}[R_t | S_t = s] \quad (5)$$

$$= \mathbb{E}[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | S_t = s] \quad (6)$$

$$= \mathbb{E}[r_{t+1} + \gamma(r_{t+2} + \gamma r_{t+3} + \dots) | S_t = s] \quad (7)$$

$$= \mathbb{E}[r_{t+1} + \gamma R_{t+1} | S_t = s] \quad (8)$$

$$= \mathbb{E}[r_{t+1} + \gamma v(S_{t+1}) | S_t = s] \quad (9)$$

Value equation decomposes into 2 terms:

- Immediate reward  $r_{t+1}$
- Discounted return of successor state  $\gamma v(S_{t+1})$

Note for the mathematically orthodox: Between the two last derivation steps we swept under the carpet that we are using the Tower rule/Law of Total expectation to replace the expected return for state with that of the successor state, which is beyond what we expect for this course.

# Forms of the Bellman Equation for MRPs

- Expectation notation:

$$v(s) = \mathbb{E}[R_t | S_t = s]$$

- Sum notation (expectation written out):

$$v(s) = \mathcal{R}_s + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'} v(s') \quad (10)$$

We have  $n$  of these equations, one for each state.

- Vector notation:

$$\mathbf{v} = \mathcal{R} + \gamma \mathcal{P} \mathbf{v} \quad (11)$$

The vector  $\mathbf{v}$  is  $n$ -dimensional.

# Direct solution

The Bellman equation is a linear, self-consistent equation:

$$\mathbf{v} = \mathcal{R} + \gamma \mathcal{P} \mathbf{v}$$

we can solve for it directly:

$$\mathbf{v} = \mathcal{R} + \gamma \mathcal{P} \mathbf{v} \tag{12}$$

$$(\mathbf{1} - \gamma \mathcal{P}) \mathbf{v} = \mathcal{R} \tag{13}$$

$$\mathbf{v} = (\mathbf{1} - \gamma \mathcal{P})^{-1} \mathcal{R} \tag{14}$$

Matrix inversion is computational expensive at  $\mathcal{O}(n^3)$  for  $n$  states (e.g. Backgammon has  $10^{20}$  states), so direct solution only feasible for small MRPs.

Fortunately there are many iterative methods for solving large MRPs:


- Dynamic programming
- Monte-Carlo evaluation
- Temporal-Difference learning

These are at the core of Reinforcement learning, we will learn all 3 algorithms.

By the way you have met the solution of **self-consistent equations** before, whenever you solved for a set of  $n$  linear equations in  $n$  unknowns you exploited that the equations and the unknowns had to be self-consistent (i.e. related to each other by the common structure of the problem).


# Myopic MRP value function: $\gamma = 0$

What is the value of being in a state?



Immediate reward and total return match, but what if  $\gamma > 0$ ?

# Far sighted MRP value function $\gamma = 1$





## Far sighted MRP value function $\gamma = 1$ II

- Check for **self-consistency**  $v(C3) = ?$

Using Bellman Equation:

$$v(C3) = -2 + 1 \times 0.6 \times 10 + 1 \times 0.4 \times 0.8 = 4.32 \approx 4.3$$

There are cycles in the graph, why is the value of some states not infinite?

This is self-consistent (numbers in figure are rounded to one digit for space reasons)

- Check for **self-consistency**  $v(C1) = ?$


Using Bellman Equation:  $v(C1) =$

$$-2 + 1 \times 0.5 \times -23 + 1 \times 0.5 \times 1.5 = -12.75 \approx -13$$

This is self-consistent (numbers in figure are rounded to one digit for space reasons)

# From state to action: policy

## Example (Coffee Process)



States are 1, 2, 3 and 4. What to do where?

$\pi(1)$  = turn left and walk


$\pi(2)$  = go straight and walk

$\pi(3)$  = turn right and walk

$\pi(4)$  = turn right and walk

# Rock, Paper & Scissors Process

## Example (Rock, Paper & Scissors Process)



Following a rigid policy can be disadvantageous and exposes the agent to being systematically exploited (see "Dutch Book" argument).

## Definition (Policy)

A policy  $\pi_t(a, s) = P[A_t = a | S_t = s]$  is the conditional probability distribution to execute an action  $a \in \mathcal{A}$  given that one is in state  $s \in \mathcal{S}$  at time  $t$ .

The general form of the policy is called a **probabilistic** or **stochastic** policy, so  $\pi$  is a probability. If for a given state  $s$  only a single  $a$  is possible, then the policy is **deterministic**:  $\pi(a, s) = 1$  and  $\pi(a', s) = 0, \forall a \neq a'$ . A shorthand is to write  $\pi_t(s) = a$ , implying that the function  $\pi$  returns an action for a given state. Now we "only" need to work out how to choose an action ...

# Lottery decision making

## Example

Optimal decision maximises our expected return

Actions	Reward
$a_1$ : play	$s_1$ : Win the lottery
$a_2$ : save	$s_2$ : Lose the lottery

$$a^* = \arg \max_{a_i} \sum_{j=1}^2 \mathcal{R}_{s_j}^{a_i} P[s_j | a_i] \quad (15)$$

Where  $a_i$  are the actions available in state  $s_i$ , i.e.

$a_i \in \mathcal{A}(s_i)$ .

$$\begin{array}{ll}
 P[s_1 | a_1] = 10^{-7} & \mathcal{R}_1^1 = 500,000 \text{ USD} \\
 P[s_2 | a_1] = 1 - 10^{-7} & \mathcal{R}_1^2 = -1 \text{ USD} \\
 P[s_1 | a_2] = 0 & \mathcal{R}_2^1 = 0 \text{ USD} \\
 P[s_2 | a_2] = 1 & \mathcal{R}_2^2 = 0 \text{ USD}
 \end{array}$$

What is the optimal action for this decision problem?